

# CS624 - Analysis of Algorithms

## Heaps

February 6, 2024

# Binary Trees





## Definition (Binary Tree)

A **binary tree** is either

- ▶ a **node** with two children, called *left* and *right*, which are also binary trees, and optionally a *data* field; or
- ▶ NIL, representing the empty tree

## Examples (Binary trees with and without data)

NIL




# Binary Trees


## Definition (Leaf Node)

A **leaf node** is a **node** whose children are both NIL.


## Examples



=



=



NIL is often omitted from tree drawings, unless the tree is empty.

# Node Height

## Definition (Height)


The **height** of a **node** in a **binary tree** is the number of edges on the longest path from the node to a **leaf node**.

The **height** of a **binary tree** is the height of its root node.

## Example

In the tree to the right, each node is labeled with its **height**.

The tree has height 4.




# Node Level

## Definition (Level)

The **level** of a **node** is the number of edges from it to the **root node**. In general, if a node has **level**  $k$ , its children both have **level**  $k + 1$ .

## Example

In the tree to the right, each node is labeled with its **level**.



There are at most  $2^k$  nodes at level  $k$ .

If the highest level is completely filled in, that level contains  $2^H$  nodes, and the tree contains  $1 + 2 + 4 + \dots + 2^H = 2^{H+1} - 1$  nodes.

## Definition (Pre-heap)


A **pre-heap** is a binary tree with the following properties:

- ▶ All **leaves** are on at most two adjacent levels.
- ▶ All **levels**, except maybe the lowest, are completely filled.
- ▶ The lowest **level** is filled without gaps, from the left.


A **pre-heap** can be efficiently, compactly represented using an array.

# Pre-Heap Representation

**Logical view**



**Efficient representation**



$$\text{Left}(k) = 2k$$

$$\text{Right}(k) = 2k + 1$$

$$\text{Parent}(k) = \lfloor k/2 \rfloor \quad \text{if } k > 1$$

Levels 0, 1, and 2 are completely filled in and contain  $2^3 - 1$  nodes. Level 3 is partly filled in from the left.

## Observation

Suppose we have a **pre-heap** with  $n$  nodes and height  $H$ .  
Then  $2^H \leq n \leq 2^{H+1} - 1 < 2^{H+1}$ . So  $H = \lfloor \log_2 n \rfloor$ .



# Pre-Heap Properties

## Lemma


*In a pre-heap with  $n$  elements, there are  $\lceil \frac{n}{2} \rceil$  leaves.*

## Proof.

- ▶ Let  $H$  be the height. All nodes on level  $H$  are leaves.
- ▶ Some of the rightmost nodes on level  $H - 1$  may be leaves.
- ▶ The parent of node  $n$  is  $\lfloor \frac{n}{2} \rfloor$ , and that node is the last node of height 1 on level  $H - 1$ , since level  $H$  is filled from the left.
- ▶ That is, all nodes after  $\lfloor \frac{n}{2} \rfloor$  on level  $H - 1$  are leaves, and all nodes on level  $H$  are leaves.
- ▶ So there are  $n - \lfloor \frac{n}{2} \rfloor = \lceil \frac{n}{2} \rceil$  leaves.



# Pre-Heap Properties



The parent of node  $n$  is  $\lfloor \frac{n}{2} \rfloor$ , and that node is the last node of **height** 1 on **level**  $H - 1$ , since **level**  $H$  is filled from the left.

# Pre-Heap Properties

## Corollary

*In a pre-heap with height  $H$ , there are at most  $2^H$  leaves.*

## Proof.

If  $n$  is the number of elements in the pre-heap, we know that

$$2^H \leq n \leq 2^{H+1} - 1 < 2^{H+1}$$

Then by the Lemma, the number of leaves is

$$\left\lceil \frac{n}{2} \right\rceil \leq \frac{2^{H+1}}{2} = 2^H$$



# Pre-Heap Properties

## Theorem

*In a pre-heap with  $n$  elements, there are at most  $\frac{n}{2^h}$  nodes at height  $h$ .*

## Proof.

- ▶ We have just seen that there are at most  $2^H$  leaves in such a tree, and the leaves are just the nodes at height 0.
- ▶ If we take away the leaves, we have a smaller pre-heap with at most  $2^{H-1}$  leaves, and these leaves are exactly the nodes at height 1 in the original tree.
- ▶ Continuing, we see that there are at most  $2^{H-h}$  nodes at height  $h$  in the original tree, therefore  $2^{H-h} = \frac{2^H}{2^h} \leq \frac{n}{2^h}$



## Definition (Heap)

A **heap** is a **pre-heap** where each node contains a key, the keys are comparable, and each node satisfies the **heap properties**:

1. The node's key is greater than or equal to the keys of its children.
2. The node's left and right subtrees are also **heaps**.



Specifically, this is called a **max-heap**; the root has the maximum key.

Another way of phrasing the **heap** conditions would be:

- ▶ The key at each node is greater than or equal to the key in any descendant of that node.

## Example: Two Heaps With the Same Set of Keys

The shape of a **heap** with  $n$  elements is uniquely determined, since it is a **pre-heap**, but the arrangement of the elements is not.




# Building a Heap

How can we (efficiently) build a heap?

**Input:** A **pre-heap** represented by an array.


**Strategy:**



# Building a Heap

How can we (efficiently) build a heap?

**Input:** A **pre-heap** represented by an array.



**Strategy:** Let's try **divide and conquer**:

- ▶ **Sub-problems:** Turn the left and right children of the into **heaps**.
- ▶ **Combine:** Given a **pre-heap** whose left and right children are **heaps**, convert the whole thing into one **heap**.



# Building a Heap, Recursively

Initial call:  $\text{BuildHeapRec}(A, 1)$ , with  $\text{Heapsize}(A) \leftarrow \text{Length}(A)$ .

---

**Algorithm 1**  $\text{BuildHeapRec}(A, i)$

---

**Ensure:** The subtree rooted at  $A[i]$  is a **heap**.

```
1: if  $i > \text{Heapsize}(A)$  then
2:   // Then  $i$  represents NIL
3:   return
4: else
5:    $l \leftarrow \text{Left}(i)$ 
6:    $r \leftarrow \text{Right}(i)$ 
7:    $\text{BuildHeapRec}(A, l)$ 
8:    $\text{BuildHeapRec}(A, r)$ 
9:    $\text{Heapify}(A, i)$ 
10: end if
```

---

# The Heapify Procedure

---

**Algorithm 2** Heapify( $A, i$ )

---

**Require:**  $1 \leq i \leq \text{Heapsize}(A)$ , and the sub-trees rooted at  $A[\text{Left}(i)]$  and  $A[\text{Right}(i)]$  (if they exist) are **heaps**.

**Ensure:** The tree rooted at  $A[i]$  is a **heap**.

```
1:  $l \leftarrow \text{Left}(i)$ 
2:  $r \leftarrow \text{Right}(i)$ 
3:  $largest \leftarrow i$ 
4: if  $l \leq \text{Heapsize}[A]$  and  $A[l] > A[i]$  then
5:    $largest \leftarrow l$ 
6: end if
7: if  $r \leq \text{Heapsize}[A]$  and  $A[r] > A[largest]$  then
8:    $largest \leftarrow r$ 
9: end if
10: if  $largest \neq i$  then
11:   exchange  $A[i] \leftrightarrow A[largest]$ 
12:   Heapify( $A, largest$ )
13: end if
```

---

# The Heapify Procedure





Comments on Heapify:

- ▶ *largest* is the index of node with the largest key, out of  $i$ , its left child (if it exists), and its right child (if it exists).

There are three cases:

1. *largest* =  $i$ . Then  $A[i]$  satisfies the **heap properties**. Done.
  2. *largest* =  $\text{Left}(i)$ . Then  $A[i]$  does not satisfy **heap property #1**.  
By exchanging  $A[i] \leftrightarrow A[\text{largest}]$ , we fix **heap property #1**, but we may have broken **heap property #2**: the left sub-tree may no longer be a heap. So we repair it by calling **Heapify** recursively. (?)
  3. *largest* =  $\text{Right}(i)$ . Similar to previous case.
- ▶ The algorithm works by letting the value  $A[i]$  “float down” to its proper position in the **heap**.

# Example: Heapify



# Running Time of Heapify

The time needed to run Heapify on a subtree of size  $n$  rooted at a given node  $i$  is (worst case)

- ▶ time  $\Theta(1)$  to fix up the relationships among the elements  $A[i]$ ,  $A[\text{Left}(i)]$ , and  $A[\text{Right}(i)]$ , plus
- ▶ time to run Heapify on a subtree rooted at one  $i$ 's children  
That subtree has size at most  $2n/3$  — the worst case occurs when the last row of the tree is exactly half full. (?)

So the running time  $T(n)$  can be characterized by the recurrence

$$T(n) \leq T(2n/3) + \Theta(1)$$

# Running Time of Heapify

The time needed to run Heapify on a subtree of size  $n$  rooted at a given node  $i$  is (worst case)

- ▶ time  $\Theta(1)$  to fix up the relationships among the elements  $A[i]$ ,  $A[\text{Left}(i)]$ , and  $A[\text{Right}(i)]$ , plus
- ▶ time to run Heapify on a subtree rooted at one  $i$ 's children  
That subtree has size at most  $2n/3$  — the worst case occurs when the last row of the tree is exactly half full. (?)

So the running time  $T(n)$  can be characterized by the recurrence

$$T(n) \leq T(2n/3) + \Theta(1)$$

This falls into case 2 of the [master theorem](#), and so we must have  $T(n) = O(\log n)$ .

# Running Time of BuildHeapRec

We can express the running time for BuildHeapRec with this recurrence:

$$T(n) \leq 2T(2n/3) + \Theta(\log n)$$


By the [master theorem](#) (case 1), that gives us  $\Theta(n^{\log_{3/2} 2}) = O(n^{1.71})$ . This recurrence is a coarse bound, though; it fails to capture the fact that the sum of the subproblem sizes is less than  $n$ .

We can change the order that we tackle the subproblems. That leads to an iterative algorithm and a better analysis.


# Changing the Order of Subproblems

The *real work* happens in the combine step: Heapify.  
What order is Heapify first called on a node?

**Recursive**




**Iterative**





# Changing the Order of Subproblems

The *real work* happens in the combine step: Heapify.  
What order is Heapify first called on a node?



**Optimization:** There's no need to call Heapify on a leaf node.

# Building a Heap

**Iterative algorithm:** The heap is built from the bottom up, starting at the first non-leaf node.

---

**Algorithm 3** BuildHeap( $A$ )

---

**Ensure:** The tree rooted at  $A[1]$  is a **heap**.

```
1: Heapsize[ $A$ ]  $\leftarrow$  Length[ $A$ ]  
2: for  $i \leftarrow \lfloor \text{Length}[A]/2 \rfloor$  to 1 do  
3:   Heapify( $A, i$ )  
4: end for
```

---

To prove that this is correct we use the following **loop invariant**:

**Lemma (Loop Invariant)**

*Let  $n = \text{Length}(A)$ . At the start of each iteration of the **for** loop, each node  $i + 1, i + 2, \dots, n$  is the root of a **heap**.*

# Proof of Correctness

## Proof.

**Initialization:** On the first iteration of the loop,  $i = \lfloor n/2 \rfloor$ . Each node  $\lfloor n/2 \rfloor + 1, \lfloor n/2 \rfloor + 2, \dots, n$  is a **leaf** and is thus the root of a trivial **heap**.

### Maintenance:

- ▶ Assume it is true for loop  $i$ : each node  $i + 1, \dots, n$  is the root of a **heap**.
- ▶ Goal: show that at the end of iteration  $i$ , the invariant is true for  $i - 1$ .
- ▶ By the LI, both children of  $i$ , namely  $2i$  and  $2i + 1$  (if they exist) are **heaps**. That satisfies the **precondition** for `Heapify`.
- ▶ The call to `Heapify(A, i)` makes  $i$  a **heap** (**postcondition**).
- ▶ Furthermore, all nodes which are not descendants of  $i$  are untouched by the call to `Heapify(A, i)`, and so we can conclude that each node  $i, i + 1, \dots, n$  is now the root of a heap.

**Termination:** The loop exits when  $i = 0$ , and the loop invariant implies that node 1 is the root of a **heap**. □

# Running Time of BuildHeap

- ▶ The number of elements of the heap at height  $h$  is  $\leq \frac{n}{2^h}$ , and the cost of running Heapify on a node of height  $h$  is  $O(h)$ .
- ▶ The root of a heap of  $n$  elements has height  $\lfloor \log_2 n \rfloor$ .
- ▶ Therefore the worst-case cost of running BuildHeap on a heap of  $n$  elements is bounded by

$$\sum_{h=0}^{\lfloor \log_2 n \rfloor} \frac{n}{2^h} O(h) = O\left(n \sum_{h=0}^{\lfloor \log_2 n \rfloor} \frac{h}{2^h}\right) = O(n)$$

Since the sum converges, so we don't care what the upper bound of the summation is.

# Heap Properties

Heaps give us partial information about the order of their elements.

- ▶ We can tell immediately what the largest element is.
- ▶ They are really cheap to build.
- ▶ They are reasonably cheap to update.
- ▶ They can be stored in a simple array.

This makes them very useful for various applications.

---

**Algorithm 4** Heapsort( $A$ )

---

**Ensure:**  $A$  is sorted

```
1: BuildHeap( $A$ )
2: for  $i \leftarrow \text{Length}[A]$  to 2 do
3:   exchange  $A[1] \leftrightarrow A[i]$ 
4:   Heapsize[ $A$ ]  $\leftarrow$  Heapsize[ $A$ ] - 1
5:   Heapify( $A$ , 1)
6: end for
```

---

The call to BuildHeap takes time  $O(n)$ . Each of the  $n - 1$  calls to Heapify takes time  $O(\log n)$ . Hence the total running time (in the worst case) is  $O(n \log n)$ .

# Priority Queues

## Definition

A **priority queue** is a data structure that maintains a set  $S$  of elements, each with an associated value called a *key*. (As usual, the keys must be comparable.)

The priority queue supports the following operations:

- ▶  $\text{Insert}(S, x)$  inserts the element  $x$  into the set  $S$ .
- ▶  $\text{Maximum}(S)$  returns the element of  $S$  with the largest key.
- ▶  $\text{ExtractMax}(S)$  removes and returns the element of  $S$  with the largest key.
- ▶  $\text{IncreaseKey}(S, x, k)$  increases the value of element  $x$ 's key to the new value  $k$ , which must be at least as large as  $x$ 's current key value.

A **priority queue** can be implemented using a **heap**.

# Priority Queue Operations

---

**Algorithm 5** HeapMaximum( $A$ )

---

**Require:**  $\text{Heapsize}(A) \geq 1$

1: **return**  $A[1]$

---

Obviously, the run time is  $O(1)$ .

---

**Algorithm 6** HeapExtractMax( $A$ )

---

**Require:**  $\text{Heapsize}(A) \geq 1$

1:  $\text{maxx} \leftarrow A[1]$

2:  $A[1] \leftarrow A[\text{Heapsize}[A]]$

3:  $\text{Heapsize}[A] \leftarrow \text{Heapsize}[A] - 1$

4: Heapify( $A, 1$ )

5: **return**  $\text{maxx}$

---

Here the running time is dominated by the call to Heapify, so it is  $O(\log n)$ .



---

**Algorithm 7** HeapIncreaseKey( $A, i, key$ )





---

**Require:**  $key \geq A[i]$

- 1:  $A[i] \leftarrow key$
  - 2: **while**  $i > 1$  **and**  $A[\text{Parent}(i)] < A[i]$  **do**
  - 3:   exchange  $A[i] \leftrightarrow A[\text{Parent}(i)]$
  - 4:    $i \leftarrow \text{Parent}(i)$
  - 5: **end while**
- 

We just increase the key of  $A[i]$ , and then let that node “float up” to its proper position.

## Example: HeapIncreaseKey



---

**Algorithm 8** HeapInsert( $A, key$ )

---

**Require:** Heapsize( $A$ ) < Length( $A$ ), or  $A$  can grow

- 1: Heapsize[ $A$ ]  $\leftarrow$  Heapsize[ $A$ ] + 1
  - 2:  $A$ [Heapsize[ $A$ ]]  $\leftarrow -\infty$
  - 3: HeapIncreaseKey( $A$ , Heapsize[ $A$ ],  $key$ )
- 

The running time here is again  $O(\log_2 n)$ .

Thus, a heap supports any priority queue operation on a set of size  $n$  in  $O(\log n)$  time.