

Medians and Order Statistics

CS 624 — Analysis of Algorithms

February 20, 2024

- ▶ **Tentative dates:** The midterm exam will take place on either March 5 (Tuesday) or March 7 (Thursday), in class.
- ▶ Covered material: Induction, runtime analysis, sorting (mergesort, insertion sort, quicksort, heapsort, lower bounds), heaps, medians and order statistics, binary search trees.
Not covered: dynamic programming.
- ▶ The previous class will be partly a review class.
- ▶ Prepare your own questions to ask me!

- ▶ Probably 4 questions. Assume every topic will be covered.
- ▶ No books, no computers, no cellphones/smartphones/tablets, strictly no friends.
- ▶ You may bring up to 20 pages of **handwritten notes**.
(That is, 20 pieces of paper, up to letter size.)
No printouts, no photocopies.

Medians and Order Statistics

Definition (Order Statistic)

The i^{th} **order statistic** is the i^{th} smallest element of a set of n elements.

In particular:

- ▶ **minimum** = 1^{st} order statistic
- ▶ **maximum** = n^{th} order statistic
- ▶ **median**: “half-way point” of the set
 - ▶ the **lower median** is at $\lfloor (n+1)/2 \rfloor$
 - ▶ the **upper median** is at $\lceil (n+1)/2 \rceil$
 - ▶ same when n is odd, different when n is even
 - ▶ for simplicity, “median” refers to the **lower median**

Selection Problem

Definition (Selection Problem)

The **selection problem** is defined as follows:

- ▶ Input: A set A of n **distinct** numbers and a number k , with $1 \leq k \leq n$.
- ▶ Output: the element $x \in A$ that is larger than exactly $k - 1$ other elements of A (that is, the k^{th} **order statistic**).

Can be solved in $O(n \log n)$ time. How?

Selection Problem

Definition (Selection Problem)

The **selection problem** is defined as follows:

- ▶ Input: A set A of n **distinct** numbers and a number k , with $1 \leq k \leq n$.
- ▶ Output: the element $x \in A$ that is larger than exactly $k - 1$ other elements of A (that is, the k^{th} **order statistic**).

Can be solved in $O(n \log n)$ time. How?

There are faster, linear-time algorithms.

- ▶ For the special cases when $k = 1$ and $k = n$.
- ▶ For the general problem.

Minimum and Maximum

The **minimum** or **maximum** can be found in $\Theta(n)$ time.


- ▶ Simply scan all the elements and find the smallest/largest.

Simultaneous Minimum and Maximum


Some applications need to determine both the **minimum** and **maximum** of a set of elements.

- ▶ Example: Graphics program trying to fit a set of points onto a rectangular display.

Calculating the **minimum** and **maximum** *independently* requires $2n - 2$ comparisons. Can we reduce this number?



Simultaneous Minimum and Maximum




The algorithm sketch:

- ▶ maintain *min* and *max* elements seen so far
- ▶ process elements in *pairs*, compare to get smaller and larger
- ▶ compare smaller to *min* and larger to *max*, update

There are 3 comparisons per pair, and $\lfloor n/2 \rfloor$ pairs.

Simultaneous Minimum and Maximum



The algorithm sketch:

- ▶ maintain min and max elements seen so far
- ▶ process elements in *pairs*, compare to get smaller and larger
- ▶ compare smaller to min and larger to max , update

There are 3 comparisons per pair, and $\lfloor n/2 \rfloor$ pairs.

Simultaneous Minimum and Maximum

Analysis:

- ▶ For odd n : initialize min and max to $A[1]$. Pair the remaining elements. So, number of pairs = $\lfloor n/2 \rfloor$.
- ▶ For even n : initialize min to the smaller of the first pair and max to the larger. So, remaining number of pairs = $(n - 2)/2 < \lfloor n/2 \rfloor$.
- ▶ Total number of comparisons, $C \leq 3\lfloor n/2 \rfloor$.
- ▶ For odd n : $C = 3\lfloor n/2 \rfloor$.
- ▶ For even n : $C = 3(n - 2)/2 + 1 = 3n/2 - 2 < 3\lfloor n/2 \rfloor$.

Finding k^{th} Smallest Element

Can we use a similar method for any order statistic in linear time?

- ▶ The cost of finding the k^{th} order statistic using either of these methods is $\Theta(kn)$. If k is fixed, this is $\Theta(n)$.
- ▶ If k is not fixed, this is not so good. For instance, suppose we want to find the **median**. Then k is $n/2$, and the cost is $\Theta(n^2)$, worse than sorting the array.

Is there an $O(n)$ time (independent of k) algorithm for selecting the k^{th} **order statistic**?

Finding k^{th} Smallest Element

Can we use a similar method for any order statistic in linear time?

- ▶ The cost of finding the k^{th} order statistic using either of these methods is $\Theta(kn)$. If k is fixed, this is $\Theta(n)$.
- ▶ If k is not fixed, this is not so good. For instance, suppose we want to find the **median**. Then k is $n/2$, and the cost is $\Theta(n^2)$, worse than sorting the array.

Is there an $O(n)$ time (independent of k) algorithm for selecting the k^{th} **order statistic**? Yes:

- ▶ a simple algorithm with expected $O(n)$ complexity
- ▶ a variant with worst-case $O(n)$ complexity

General Selection Problem

Given: array A of size n and k such that $1 \leq k \leq n$

- ▶ If the array A were **sorted**, we would simply find the k^{th} **order statistic** at $A[k]$. But we don't actually care if A is **completely sorted**, as long as $A[k]$ contains the right element.
- ▶ That is one of the properties that Quicksort establishes:
*Once Partition chooses a **pivot** and that call to Partition completes, that **pivot** never moves again.*
- ▶ We modify Quicksort to eliminate unnecessary work:
We only recur on the side containing k .
- ▶ In the average case, the cost of the Partition steps should be

$$n + \frac{n}{2} + \frac{n}{4} + \frac{n}{8} + \dots = 2n$$

That is, $O(n)$ average-case complexity.


Randomized Select

Algorithm 1 RandomizedSelect(A, p, r, k_{rel})

Require: $1 \leq k_{\text{rel}} \leq r - p + 1$

```
1: if  $p = r$  then  
2:   return  $A[p]$   
3: end if  
4:  $q \leftarrow \text{RandomizedPartition}(A, p, r)$   
5:  $q_{\text{rel}} \leftarrow q - p + 1$   
6: if  $k_{\text{rel}} = q_{\text{rel}}$  then  
7:   return  $A[q]$   
8: else if  $k_{\text{rel}} < q_{\text{rel}}$  then  
9:   return RandomizedSelect( $A, p, q - 1, k_{\text{rel}}$ )  
10: else  
11:   return RandomizedSelect( $A, q + 1, r, k_{\text{rel}} - q_{\text{rel}}$ )  
12: end if
```

Randomized Select



Notation used in the algorithm RandomizedSelect:

- ▶ p , q , and r are indices in the original array A .
- ▶ q_{rel} is the 1-based index of the pivot $A[q]$ in the subarray $A[p \dots r]$ — that is, *relative* to the range $[p .. r]$.

We see that Randomized-Select is divided into 3 cases:

1. $q_{\text{rel}} < k_{\text{rel}}$, so we search for the $(k_{\text{rel}} - q_{\text{rel}})^{\text{th}}$ element in $A[q + 1 .. r]$
2. $q_{\text{rel}} = k_{\text{rel}}$, so we found it, and we return $A[q]$
3. $q_{\text{rel}} > k_{\text{rel}}$, so we search for the $k_{\text{rel}}^{\text{th}}$ element in $A[p .. q - 1]$

Worst-case complexity:

- ▶ $\Theta(n^2)$ — (Like Quicksort) We could get unlucky and always recur on a subarray that is only one element smaller.

Average-case complexity:

- ▶ $\Theta(n)$ — Intuition: Because the pivot is chosen at random, we expect that we get rid of half of the list each time we choose a random pivot q .
- ▶ Why $\Theta(n)$ and not $\Theta(n \log n)$?

Average-Case Analysis

Let $C(n, i)$ denote the average running time of `RandomizedSelect`($A, 1, n, i$).

Let $T(n)$ denote the **worst average-case** time of computing *any* i^{th} element of an array of size n using `RandomizedSelect`. That is:

$$T(n) = \max \{C(n, i) \mid 1 \leq i \leq n\}$$

We will prove that $T(n) = O(n)$.

Average-Case Analysis

The cost of Partition is $O(n)$, so we can bound it by an for some a .

Therefore:

$$C(n, i) \leq an + \frac{1}{n} \left(\sum_{q=1}^{i-1} C(n-q, i-q) + 0 + \sum_{q=i+1}^n C(q-1, i) \right)$$

- ▶ The call to RandomizedSelect has two parts:
 - ▶ the Partition call, whose cost is an , and
 - ▶ the recursive call, whose cost varies depending on the location of the pivot which we denote q (really should be q_{rel}).
- ▶ We assume that the pivot is equally likely to wind up in any of the n positions in the array, and we average over all those n possibilities.
- ▶ Inside the parentheses is the sum of the n possible pivots q :
 - ▶ the first term is if the pivot falls before i
 - ▶ the second term is if the pivot is exactly i (we just return)
 - ▶ the third term is if the pivot falls after i

Average-Case Analysis

$$\begin{aligned}C(n, i) &\leq an + \frac{1}{n} \left(\sum_{q=1}^{i-1} C(n-q, i-q) + 0 + \sum_{q=i+1}^n C(q-1, i) \right) \\&\leq an + \frac{1}{n} \left(\sum_{q=1}^{i-1} T(n-q) + \sum_{q=i+1}^n T(q-1) \right) \\&\leq \max \left\{ an + \frac{1}{n} \left(\sum_{q=1}^{i-1} T(n-q) + \sum_{q=i+1}^n T(q-1) \right) \mid 1 \leq i \leq n \right\} \\&= an + \max \left\{ \frac{1}{n} \left(\sum_{q=1}^{i-1} T(n-q) + \sum_{q=i+1}^n T(q-1) \right) \mid 1 \leq i \leq n \right\}\end{aligned}$$

Average-Case Analysis — Explanation

- ▶ Note: i is a separate variable, different from the i in the left-hand side $C(n, i)$. In fact, in the final inequality for $C(n, i)$, the right-hand side no longer depends on i .
- ▶ Substituting in the definition of $T(n)$, we get:

$$\begin{aligned} T(n) &= \max \{C(n, i) \mid 1 \leq i \leq n\} \\ &\leq an + \max \left\{ \frac{1}{n} \left(\sum_{q=1}^{i-1} T(n-q) + \sum_{q=i+1}^n T(q-1) \right) \mid 1 \leq i \leq n \right\} \end{aligned}$$

We'll **guess** that $T(n) = O(n)$ and **prove** by induction that $T(n) = Cn$ satisfies the inequality above.

Average-Case Analysis — Proof by Induction

Theorem

Suppose that

$$T(n) \leq an + \max \left\{ \frac{1}{n} \left(\sum_{q=1}^{i-1} T(n-q) + \sum_{q=i+1}^n T(q-1) \right) \mid 1 \leq i \leq n \right\}$$

Then $T(n) \leq Cn$ for some $C > 0$.

Proof.

Base Case: We can arrange that this is true for $n = 2$ by making sure (when we finally figure out an appropriate value for C) that $C \geq a$.

Inductive Case: We must show that $T(n) \leq Cn$.

IH: Assume that $T(k) \leq Ck$ for all $1 \leq k < n$.

Average-Case Analysis — Proof by Induction

Proof.

We start with the recursive inequality:

$$\begin{aligned} T(n) &\leq an + \max \left\{ \frac{1}{n} \left(\sum_{q=1}^{i-1} T(n-q) + \sum_{q=i+1}^n T(q-1) \right) \mid 1 \leq i \leq n \right\} \\ &\leq an + \max \left\{ \frac{C}{n} \left(\sum_{q=1}^{i-1} (n-q) + \sum_{q=i+1}^n (q-1) \right) \mid 1 \leq i \leq n \right\} \quad (\text{by IH}) \\ &= an + \max \left\{ \frac{C}{n} \left((i-1)n - \frac{(i-1)i}{2} + \frac{(n-1)n}{2} - \frac{(i-1)i}{2} \right) \mid 1 \leq i \leq n \right\} \\ &= an + \max \left\{ \frac{C}{n} \left((i-1)n - (i-1)i + \frac{(n-1)n}{2} \right) \mid 1 \leq i \leq n \right\} \end{aligned}$$

Average-Case Analysis — Proof by Induction

Proof.

- ▶ We have to find the maximum of $(i - 1)n - (i - 1)i = -i^2 + (n + 1)i - n$ between $i = 1$ and $i = n$. This is a concave function of i ; in fact, it's an “upside-down parabola”, and so its maximum occurs where the derivative is 0.
- ▶ The derivative is $-2i + (n + 1)$ and this is 0 when $i = \frac{n+1}{2}$.
- ▶ So the maximum value of the expression $(i - 1)n - (i - 1)i$, which is also $(i - 1)(n - i)$, is

$$\left(\frac{n+1}{2} - 1\right)\left(n - \frac{n+1}{2}\right) = \frac{n-1}{2} \frac{n-1}{2} = \frac{(n-1)^2}{4}$$

Average-Case Analysis — Proof by Induction

Proof.

So we have

$$\begin{aligned}T(n) &\leq an + \frac{C}{n} \left(\frac{(n-1)^2}{4} + \frac{(n-1)n}{2} \right) \\&= an + \frac{C}{n} \left(\frac{n^2 - 2n + 1}{4} + \frac{n^2 - n}{2} \right) \\&= an + \frac{C}{n} \left(\frac{3n^2}{4} - n + \frac{1}{4} \right) \\&= an + C \left(\frac{3n}{4} - 1 + \frac{1}{4n} \right) \\&\leq an + C \frac{3n}{4} \quad \text{for } n \geq 1 \\&= \left(a + \frac{3}{4}C \right) n\end{aligned}$$

Average-Case Analysis — Proof by Induction

Proof.

So we can fix C finally so that

- ▶ $C \geq a$, and
- ▶ $a + (3/4)C \leq C$

For instance, $C = 4a$ would work.

Then we get $T(n) \leq Cn$ and we are done.





Selection in Worst-Case Linear-Time

The previous algorithm has *expected* $O(n)$ running time, but the worst case is $O(n^2)$, like Quicksort.

There is a variant that runs in $O(n)$ time in the worst case:

- ▶ instead of picking a random **pivot**, use the **median of medians**
- ▶ divide the input range $A[p .. r]$ into $\lfloor n/5 \rfloor$ groups of 5
- ▶ sort each group to find its **median**
- ▶ recursively find the **median** of the group **medians**
- ▶ that is a good enough **pivot** to guarantee linear complexity





Let m be the **median of medians**. In each *full* column to the left or right of m , 3 of the 5 elements are $<$ or $>$ to m , respectively.

Analysis of Select

Let $T(n)$ be the running time of Select using **median of medians**:

- ▶ sort each group of 5 to find its **median** — $O(5^2)\lfloor n/5 \rfloor = \Theta(n)$
- ▶ recursively find the **median** of the group **medians** — $T(\lfloor n/5 \rfloor)$
- ▶ partition using the **median of medians** as **pivot** — $\Theta(n)$
- ▶ recur on one of the partitions — $\leq T(7n/10)$ (!?)

To summarize:

$$T(n) \leq T(n/5) + T(7n/10) + \Theta(n)$$

We can use **guess and prove** to show that this is $O(n)$.

We also know $T(n) = \Omega(n)$, so we get $T(n) = \Theta(n)$.