

Amortized Analysis

CS 624 — Analysis of Algorithms

April 9, 2024



Amortized Analysis

- ▶ **Amortized analysis** is a technique for analyzing the efficiency of an algorithm in situations where the cost of a *single* operation can vary greatly.
- ▶ An **amortized analysis** is any strategy for analyzing a *sequence of operations* to show that the *average* cost of an operation is small, even though a single operation within the sequence might be relatively expensive.
- ▶ We'll explore three “methods” of **amortized analysis**:
 - ▶ the **aggregate method**
 - ▶ the **accounting method**
 - ▶ the **potential method**

Incrementing a Binary Counter

- ▶ Suppose we have a binary counter: an array of k bits, where each bit can be flipped independently of the others.
The cost of flipping a bit is 1, and that there are no other costs to consider.
- ▶ There is only one public operation: INCREMENT.
This is a very simple example, since there is only one kind of legal sequence of operations.
- ▶ What happens when we start from 0 (that is, all the bits are 0) and start counting?

Incrementing a Binary Counter

counter value	bit index								incremental cost	total cost
	7	6	5	4	3	2	1	0		
0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	1	1	1
2	0	0	0	0	0	0	1	0	2	3
3	0	0	0	0	0	0	1	1	1	4
4	0	0	0	0	0	1	0	0	3	7
5	0	0	0	0	0	1	0	1	1	8
6	0	0	0	0	0	1	1	0	2	10
7	0	0	0	0	0	1	1	1	1	11
8	0	0	0	0	1	0	0	0	4	15
9	0	0	0	0	1	0	0	1	1	16
10	0	0	0	0	1	0	1	0	2	18
11	0	0	0	0	1	0	1	1	1	19
12	0	0	0	0	1	1	0	0	3	22
13	0	0	0	0	1	1	0	1	1	23
14	0	0	0	0	1	1	1	0	2	25
15	0	0	0	0	1	1	1	1	1	26
16	0	0	0	1	0	0	0	0	5	31

How expensive is INCREMENT? Is that a fair (or useful) analysis?

Running Time Analysis

Per-operation analysis:

- ▶ It's clear that the cost of any one step can be large.
- ▶ For instance, the cost of the 16th step is 5, and in fact the cost of step number 2^n is going to be $n + 1$.

On the other hand, the *typical* cost is low, and it appears that the *average* cost is low.

- ▶ For instance, the total cost of the first 16 steps is 31, so the average cost of each step (at least out to 16 steps) is less than 2.

The per-operation analysis is **unduly pessimistic**.

The Aggregate Method

We can count how many bit flips there are in the first n steps:

- ▶ Every step flips bit 0, so that contributes n bit flips.
- ▶ Every other step flips bit 1, so that contributes $\lfloor \frac{n}{2} \rfloor$ bit flips.
- ▶ Every fourth step flips bit 2, so that contributes $\lfloor \frac{n}{4} \rfloor$ bit flips etc.

The total number of bit flips in the first n steps is therefore

$$\sum_{i=0}^k \left\lfloor \frac{n}{2^i} \right\rfloor < n \sum_{i=0}^{\infty} \frac{1}{2^i} = 2n$$

No matter how far out we go, the average cost of a step is bounded above by 2. We say that the **amortized cost** of each step is ≤ 2 .

Note that the result has nothing to do with the value of k .

Amortized Analysis Methods

- ▶ The **aggregate method** analyzes *sequences of operations* to find their total **actual cost**, then averages of the number of operations to get the **amortized cost**.
- ▶ The **accounting method** requires assigning each operation a **nominal cost**, which might differ from its **actual cost**. The difference is “stored” as a “credit”. If it is proven that no sequence of operations leads to negative credit, then the **nominal cost** can be considered the **amortized cost** for each operation.
- ▶ The **potential method** is similar to the **accounting method**, but stored credit is computed from a potential function on the data structure state rather than the history of operations.

The Accounting Method

- ▶ Each operation is assigned a **nominal cost**, which may be different from its **actual cost**.
- ▶ If the **nominal cost** is greater than the **actual cost**, then we put the “extra money” aside as a **credit** to be used later.
If the **nominal cost** is less than the **actual cost**, we must use up stored **credit**.
- ▶ If we ever have negative **credit**, then the analysis is invalidated. (We must show this never happens in order to trust the analysis.)
Let c_i be the **actual cost** for the i^{th} operation, and let \hat{c}_i be the **nominal cost**. Then at each n , we must have
$$\sum_{i=1}^n \hat{c}_i \geq \sum_{i=1}^n c_i.$$
- ▶ If the accounting works out, then the **nominal cost** is considered the **amortized cost**.

The Accounting Method

- ▶ At step 1 we certainly must have $\hat{c}_1 \geq c_1$. If $\hat{c}_1 > c_1$, then we have some “money left over”.
- ▶ That means that \hat{c}_2 could actually be $< c_2$, provided that if we add in the “money left over” (i.e., the credit) from step 1, that we get a value $\geq c_2$. And so on.
- ▶ For the binary counter, we will charge an amortized cost of 2 dollars to set a bit to 1.
- ▶ When we set a bit to 1, one of those two dollars pays for the cost of setting the bit.
- ▶ The other dollar is money left over.
- ▶ We leave it with that bit, to be used later when the bit needs to be reset to 0.
- ▶ Thus at every step, each bit that is set to 1 has a dollar sitting on it.

The Accounting Method

- ▶ At each step in our process we walk to the left until we find the first 0 bit.
- ▶ We set that to 1 and reset all the bits to its right to 0.
- ▶ We charge 2 to set the new bit to 1.
- ▶ One of the dollars pays for the cost of setting the bit to 1, and the other dollar is left on the bit for later (for its reset to 0).
- ▶ All the bits to the right already have dollars left with them, so it doesn't cost anything additional to reset them to 0. That cost has already been paid for.
- ▶ Thus each step in the process is charged exactly 2 dollars and this suffices to pay for every operation, even though some operations set many bits.
- ▶ Again we see that the amortized cost of each operation is 2 (rather, ≤ 2).

The Potential Method

- ▶ Quite similar to the accounting method, but instead of associating the extra credits with individual objects in the data structure, the total amount of extra credits accumulated at any point is thought of as a “potential” (like potential energy in physics) and is associated with the data structure as a whole.
- ▶ The key thing here, though is that the potential is actually a function of the data structure itself, not how it was produced.
- ▶ In other words, if there were two different sequences of operations that could lead to the same state, the potential of that state would still be the same.

The Potential Method

- ▶ We assume that we have a potential function Φ on states of a data structure that takes a particular state D_i and produces a real number $\Phi(D_i)$.
- ▶ Our convention is that D_0 is always the initial state of the data structure.
- ▶ The amortized cost \hat{c}_i of the i^{th} operation is
$$\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1})$$
- ▶ Therefore we get

$$\sum_{i=1}^n \hat{c}_i = \sum_{i=1}^n (c_i + \Phi(D_i) - \Phi(D_{i-1})) = \sum_{i=1}^n c_i + \Phi(D_n) - \Phi(D_0)$$

The Potential Method

- ▶ To make this work we must have $\Phi(D_n) \geq \Phi(D_0)$ for all n .
- ▶ The trick is to come up with a potential function Φ that satisfies that constraint.
- ▶ In fact, usually, we just set $\Phi(D_0) = 0$.
- ▶ For the binary counter problem, D_i is simply the state of the counter after the i^{th} operation.
- ▶ Suppose we set

$$\Phi(D_i) = b_i = \text{the number of 1's after the } i^{th} \text{ operation}$$

- ▶ Thus $\Phi(D_0) = 0$, and clearly $\Phi(D_n) \geq \Phi(D_0)$ for all $n \geq 0$. This is a legal potential function.

The Potential Method

- ▶ Let us set t_i = the number of 1's that are reset during the i^{th} operation. Now the actual cost c_i of the i^{th} operation is $t_i + 1$ (because t_i bits are reset to 0, and one bit is set to 1). Also:
- ▶ If $b_i = 0$, then the i^{th} operation must have reset all the bits (there are k of them), and so $b_{i-1} = t_i = k$.
- ▶ If $b_i > 0$, then $b_i = b_{i-1} - t_i + 1$
- ▶ So in either case $b_i \leq b_{i-1} - t_i + 1$

The Potential Method

- ▶ The potential difference as we change from state D_{i-1} to state D_i is then

$$\Phi(D_i) - \Phi(D_{i-1}) \leq (b_{i-1} - t_i + 1) - b_{i-1} = 1 - t_i$$

- ▶ The amortized cost is
$$\hat{c}_i = c_i + \Phi(D_i) - \Phi(D_{i-1}) \leq (t_i + 1) + (1 - t_i) = 2$$
- ▶ We see in yet a third way that the amortized cost of each operation is ≤ 2 .

Dynamic Tables

- ▶ Suppose we have a table, like an array but not fixed-size. Like Java's `java.util.ArrayList`, C++'s `std::vector`, etc.
- ▶ We initially allocate an array of some fixed size, but after many additions the array may later become full.
We then create a larger array, copy the contents of the original into the new one, and then deallocate the original.
- ▶ Assume for simplicity that
 - ▶ Items are only inserted into the table, but never deleted.
 - ▶ When the underlying array is reallocated, the capacity doubles.
 - ▶ The table is initially empty.

What is the cost of managing a table like this?

- ▶ The cost of a single insertion into the underlying array is $\Theta(1)$.
- ▶ The costs of allocating a new array and deallocating the old one are small and fixed — $\Theta(1)$.
- ▶ The cost of copying m elements from the old array to the new one is $\Theta(m)$. (Not incurred on every insertion.)

Most of the time the cost of an insertion is very small but occasionally it is much bigger.

The Aggregate Method

The cost of the i^{th} operation is

$$c_i = \begin{cases} i & \text{if } i - 1 \text{ is an exact power of 2} \\ 1 & \text{otherwise} \end{cases}$$

Thus, the total cost of n INSERT operations is

$$\begin{aligned} \sum_{i=1}^n c_i &\leq n + \sum_{j=0}^{\lfloor \lg n \rfloor} 2^j \\ &< n + 2n \\ &= 3n \end{aligned}$$

So the amortized cost of each operation is at most 3.

The Accounting Method

The amortized cost c_i for the i^{th} insertion is 3 dollars, and this works out as follows, intuitively at least:

- ▶ \$1 pays for inserting the element itself
- ▶ \$1 is stored to move the element later when the table is doubled
- ▶ \$1 is stored to move an element in the table that was already moved from a previous table

Suppose the size of the table is m immediately after expansion. The number of elements in the table is $m/2$.

If we charge 3 dollars for each insertion, then by the time the table is filled up again, we will have $2(m/2)$ extra dollars, which pays for moving all the elements in the table to the new table.

So again we see that the amortized cost per operation is ≤ 3 .

The Potential Method

A potential function Φ tells us how much money we saved up.

- ▶ We want it to grow as elements are inserted so that when we need to double the table we have enough saved up to do it.
- ▶ Just before an expansion, we have $num[T] = size[T]$, so we need $\Phi(T) \geq num[T]$, which is exactly what costs to move every element in T into the new table.

Let Φ be 0 initially and also just after each table doubling.

The obvious function then is $\Phi(T) = 2 \cdot num[T] - size[T]$.

- ▶ Clearly $\Phi(T) \geq 0 = \Phi(empty\ table)$, so Φ is a legal potential function.

The Potential Method

Let the following subscripted variable be defined:

- ▶ num_i be the number of elements in the table *after* the i^{th} operation
- ▶ $size_i$ be the size of the table *after* the i^{th} operation
- ▶ Φ_i be the **potential** *after* the i^{th} operation
- ▶ c_i be the **actual cost** of the i^{th} operation

If the i^{th} INSERT operation does not trigger an expansion, then we have $size_i = size_{i-1}$, and the **amortized cost** of the operation, \hat{c}_i is:

$$\begin{aligned}\hat{c}_i &= c_i + \Phi_i - \Phi_{i-1} \\ &= 1 + (2 \cdot num_i - size_i) - (2 \cdot num_{i-1} - size_{i-1}) \\ &= 1 + (2 \cdot num_i - size_i) - (2(num_i - 1) - size_i) = 3\end{aligned}$$

The Potential Method

If the i^{th} INSERT operation *does* trigger an expansion, then it must be that $size_{i-1} = num_{i-1}$, so

$$size_i = 2 \cdot size_{i-1} = 2 \cdot num_{i-1} = 2 \cdot (num_i - 1)$$

Then \hat{c}_i is

$$\begin{aligned}\hat{c}_i &= c_i + \Phi_i - \Phi_{i-1} \\ &= num_i + (2 \cdot num_i - size_i) - (2 \cdot num_{i-1} - size_{i-1}) \\ &= num_i + (2 \cdot num_i - 2 \cdot (num_i - 1)) \\ &\quad - (2(num_i - 1) - (num_i - 1)) \\ &= num_i + 2 - (num_i - 1) = 3\end{aligned}$$

so again we see that the **amortized cost** of each insertion is ≤ 3 .